

ProSO

Operating Systems Project (PROSO)
Barcelona School of Informatics (FIB)
Universitat Politècnica de Catalunya (UPC)
2010-2011



Goals

- Get a good O.S. internal understanding
 - P1: Be able to program low level O.S. code, developing from scratch the basic system components
 - P2: Be able to add functionalities to an existing O.S.



Project

- Project 1
 - Implement a minimal O.S. kernel (ZeOS) [Linux based.](#)
- Project 2
 - Implement a Linux [Device Driver](#) as a Linux Kernel Module (LKM).

Course Description

- 10h/week of work
 - Lecture classes (1-2 h. not every week)
 - Lab classes
 - Scheduling-fixed (4h/week)
 - » with teacher assistance (2h/week)
 - On your own (5h/week)
- Intermediate meeting and code delivery to your advisor to follow the project evolution
- Teams of two students
 - Use the forum to put your group
 - Name; DNI; Lab group to assist

Grading

- Project 1 (Zeos) → 70%
 - Intermediate delivery P1.1 → 10%
 - Intermediate delivery P1.2 → 25%
 - Final Project delivery & Global evaluation P1.3→ (25% + 10%)
 - Questionnaire to check out the minimum concepts
- Project 2 (Linux modules)→ 30%
 - Final Project delivery → 30%
 - Questionnaire to check out the minimum concepts

Deliveries

- No one delivery is mandatory
- Intermediate deliveries
 - Evaluation & Feedback
 - You must meet with your advisor
- Important dates (Fridays)
 - P1.1: October 1st
 - P1.2: October 29th
 - P1.3: November 11th

 - P2.1: December 17th

Validation Test (Questionnaire)

- Questions about delivery (individual test)
- Important dates
 - QE1.1: October 11th
 - QE1.2: November 15th
 - QE1.3: December 13th
 - QE2.1: December 22th

Initial scheduling (can be modified)

	Monday	Tuesday	Wednesday	Thursday	Friday	
9/13/2010	T-S1.1(10/40/MAS)					9/17/2010
9/20/2010	T-S1.1(10/40/MAS)					9/24/2010
9/27/2010					S1.1	10/1/2010
10/4/2010	T-S1.2(10/40/MAS)					10/8/2010
10/11/2010	T-S1.2(10/40/MAS) QE1.1					10/15/2010
10/18/2010						10/22/2010
10/25/2010	T-S1.3(10/40/MAS)				S1.2	10/29/2010
11/1/2010						11/5/2010
11/8/2010	T-S1.3(10/40/MAS)					11/12/2010
11/15/2010	QE1.2					11/19/2010
11/22/2010					S1.3	11/26/2010
11/29/2010	T-S2.1(10/40/MAS)					12/3/2010
12/6/2010						12/10/2010
12/13/2010	QE1.3				S2.1	12/17/2010
12/20/2010			(Monday) Q2.1	(Tuesday)		12/24/2010

Holidays
Change school day
Submissions



Material


- <http://studies.ac.upc.es/FIB/PROSO>
- **RACO**
 - Deliveries
 - Forum
 - Noticeboard
- **Main source of documentation**
 - PROSO web page
 - P1 Documentation (in web page)

Requirements

- OS concepts (SO)
- Computer Architecture Concepts (EC2)
- Data Structures and Algorithms (PRED)

Overview of P1

Operating Systems Project (PROSO)
Barcelona School of Informatics (FIB)
Technical University of Catalunya (UPC)



Bibliography

- Understanding The Linux Kernel
 - Appendix A: System Startup
 - Chap.1, chap.2, chap.4
- Basic OS books
 - Galvin&Silberschatz&Peterson
 - Tanenbaum
 - ...



Goals and efforts in the Project

- OS job
- OS code and data characterization
 - Event-driven code:
 - How to manage functions
 - Independent/dependent layers
 - System initialization
 - Resource management:
 - Object representation
 - Object management: LISTS

To Take into account... (I)

- kernel vs. user-space applications
(From Robert Love "*Linux Kernel Development*")
 - The kernel does not have access to the C library
 - The kernel is coded in GNU C
 - The kernel lacks memory protection like user-space
 - The kernel cannot easily use floating point
 - The kernel has a small fixed-size stack
 - Synchronization and concurrency are concerns within the kernel
 - Portability is important

To Take into account... (II)

- Efficiency is important
 - Code that executes a lot of times
- OS has to guaranty the machine integrity
 - Robustness is important
 - User code is not reliable
- Clarity, scalability and modularity are important

To Take into account... (III)

- Some OS code has to be machine dependent
 - Assembly language
 - This code has to observe the compiler ABI
 - Application Binary Interface (ABI): Set of conventions that allows a linker to combine separately compiled modules into one unit without recompilation, such as
 - Calling conventions
 - Machine interface
 - Operating System interface

P1 files: What do you have?

- P1 documentation → see web page
- Basic Zeos files
 - Basic file's structure
 - Makefile
 - Assembler Macros
 - Basic Boot till `system.c`
 - Basic memory initialization
 - Segmentation
 - Pagination
 - Functions for accessing hardware
 - IDT (Interrupt Descriptor Table)
 - Memory
 - ...

What are you going to implement?

- Mechanisms for entering the system (P1.1)
 - Interruptions, exceptions, system calls
- Process Management (P1.2)
- Input/Output Management (P1.3)


Previous concepts you will need

- Interrupts/Exceptions management:
 - Interrupt
 - Trap
 - System gate. Handlers.
- Changing to Protected mode
- IRQ's Initialization
- Components
 - Global Descriptor Table (GDT) (Chap.4)
 - Interrupt Descriptor Table (IDT) (Chap.4)
 - Task State Segment (TSS) (Chap.4)

Working environment

- Remote boot using REMBO
 - Ubuntu 6.06
- Work in the local PC
 1. Get files from albanta to the PC (sftp)
 2. WORK
 3. **Put files** from PC to albanta (sftp)

WORK



- Execute options
 - Boot Zeos from a floppy
 - Slowdown the development process
 - **Boot Zeos on a machine emulator: Bochs**

Tools

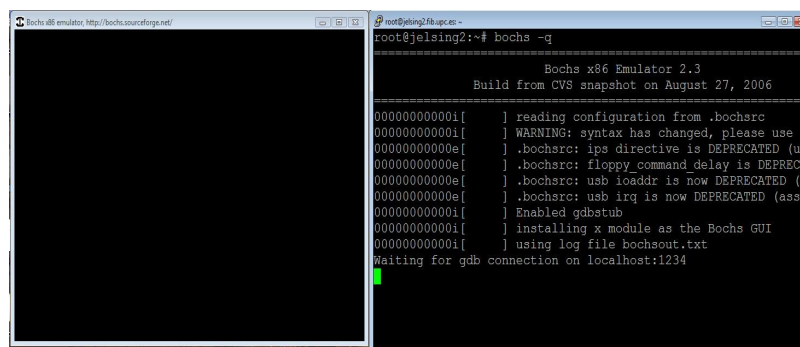
- Makefile (DONE in SO courses)
- Bochs emulator → see the documentation
- nm
 - To see where data is mapped
 - Read man pages
- objdump
 - To see **code memory addresses**
 - Read man pages

What do we ask you for in P0?

- Be familiar with the working environment
 - Linux (Development framework)
 - Bochs (emulator)
 - ZeOS (Basic files)
- Understand the boot mechanism: steps
- Build a ZeOS version from the basic files provided
- Be familiar with the debugger:
 - GDB
 - Bochs' internal debugger

Bochs emulator (v2.3)

- <http://bochs.sourceforge.net>



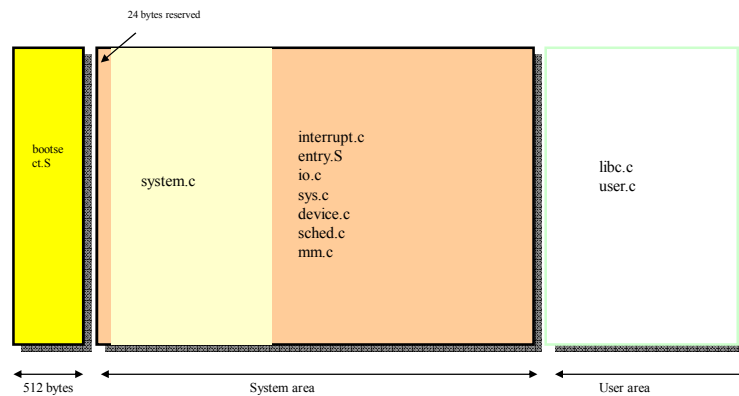
Boot Process General Overview

- Bootstrap: initializing the machine and the operating system
- Highly dependent on the computer architecture
- Intel Boot process
 - BIOS (*Basic Input/Output System*) code
 - Executed in real mode
 - Tests on the computer hardware
 - Initializes hardware devices
 - Finds the suitable boot device and loads the OS boot code (*boot loader*)
 - The “boot loader” completes the system load

ZeOS Boot file

- Created by `build.c`
 - Boot Sector (`bootsect.S`)
 - 512 bytes to fit in a floppy sector
 - System code
 - ZeOS code
 - Loads user code (start address and size must be indicated)
 - User code
 - ZeOS doesn't have a program loader

ZeOS Boot file scheme



Main() (in system.c)

- Initializes a minimum execution context
 - Prepares Null process environment
 - Initializes Segment Registers (memory management)
 - Initializes Null process kernel stack
 - Sets gdt and idtr from GDT and IDT address
- Initializes IDT (set_idt)
 - There aren't any defined interrupts
- Fills the first memory page with the system parameters
- Jumps to user space (main() function in user.c)

TO DO

- Some user code modifications
 - To get used to the working environment
- System code modifications
 - Modify `putc` to consider `\n` characters
 - Include some additional messages when booting
 - Output in the Ubuntu console
- Understand the basic ZeOSFiles.

Deliverable P1.1 **Mechanisms for entering the system**

Operating Systems Project (PROSO)
Barcelona School of Informatics (FIB)
Technical University of Catalunya (UPC)

Contents

- Mechanisms for entering the system
- Overview
 - 1. Initialization
 - 2. Management
- Procedure for entering/exiting the system
- Exceptions
- Interruptions: clock and keyboard
- System calls: write()

Mechanisms for entering the system

- Always through the IDT (Interrupt Descriptor Table)
 - From 0 to 31
 - Exceptions
 - Synchronous, produced by the CPU control unit, only after terminating the execution of an instruction
 - Non-masked interrupts
 - From 32 to 47
 - Masked Interrupts
 - Asynchronous, generated by other hardware devices at arbitrary times
 - From 48 to 255
 - Software interrupts (Traps)
 - Synchronous, explicitly requested by the application

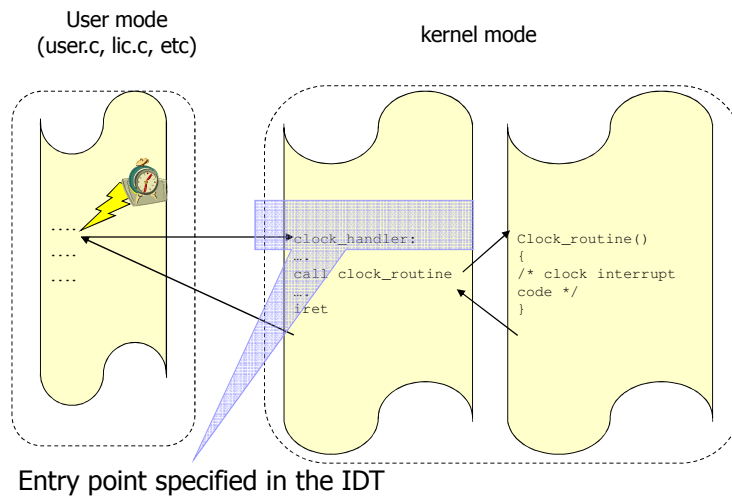
Overview: **initialization**

- Initialize the corresponding entry in the IDT
 - Each entry in the table has...
 - Interrupt number
 - Address to jump (*entry point*)
 - Privilege level
- Unmask the corresponding interrupt **if needed**
 - See enable_int function

Overview: **management code**

- Entry point → handler (entry.S)
 - Assembly code
 - Basic **hardware context** management
- Service routine (depends on the interrupt)
 - C code
 - Specific algorithm

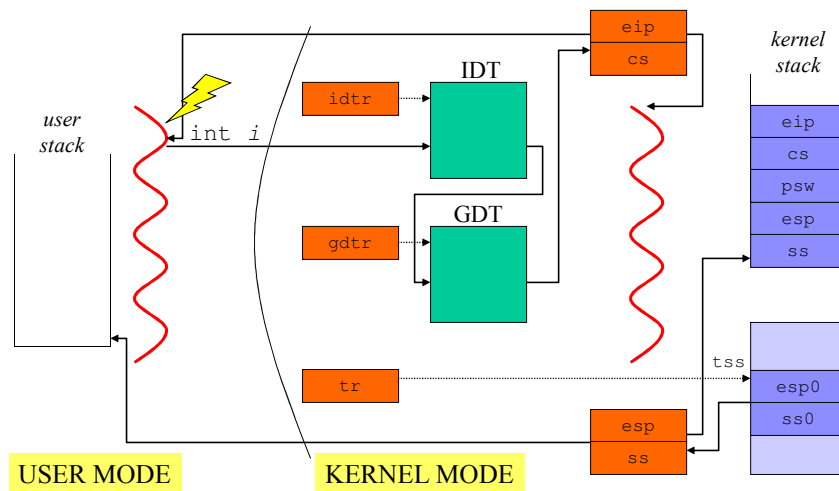
Example: clock interrupt behavior



Procedure for entering the system

- Switch to protected execution mode (*HW*)
 - User mode → Kernel mode
- Save hardware context: CPU registers
 - ss, esp, psw, cs i eip (*HW*)
 - general purpose registers (SAVE_ALL macro)
- Execute service routine (*handler*)


Procedure for entering the system



Procedure for exiting the system

- Restore hardware context
 - general purpose registers (RESTORE_ALL macro)
 - ss, esp, psw, cs i eip (iret instruction)
- Switch execution mode
 - Kernel mode -> User mode (iret instruction)

Procedure for exiting the system



These must be the contents of the kernel stack just before executing `iret` instruction


`iret` instruction empties the kernel stack and switches execution mode

Exceptions

Exceptions list

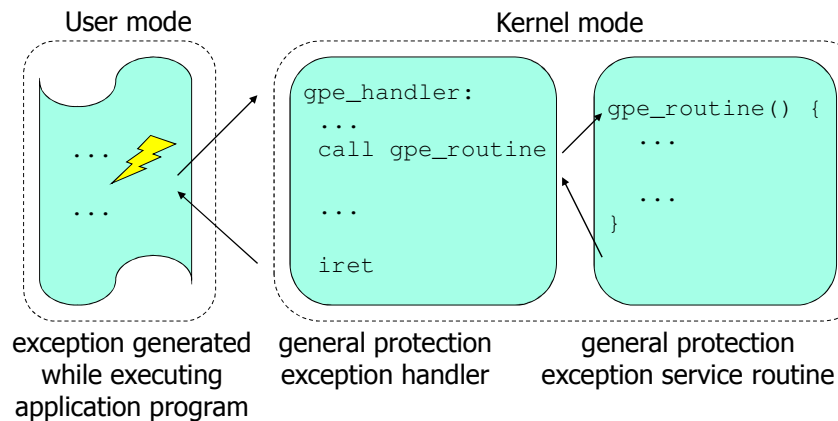
- | | |
|-----------------------------------|--------------------------|
| 0. Divide error | 10. Invalid TSS* |
| 1. Debug | 11. Segment not present* |
| 2. Not used | 12. Stack segment* |
| 3. Breakpoint | 13. General protection* |
| 4. Overflow | 14. Page fault* |
| 5. Bounds check | 15. Reserved |
| 6. Invalid opcode | 16. Floating point error |
| 7. Device not available | 17. Alignment check* |
| 8. Double fault* | 18 to 31. Reserved |
| 9. Coprocessor segment
overrun | |
- * hardware error code (4 bytes)

Handling exceptions: stack layout



For some exceptions, a hardware error code (or parameter), 4 bytes, is pushed in the kernel stack just after entering in kernel mode

Handling exceptions: overview



Handling exceptions: initialization

- Init IDT
 - one IDT entry per exception
 - void setInterruptHandler (int n, void (*f)(), int DPL)
 - n: nth IDT entry
 - f: exception handler routine address
 - DPL: max privilege level
 - e.g. `setInterruptHandler(0, divide_error_handler, KERNEL_LEVEL)`
DEFINE KERNEL_LEVEL AND USER_LEVEL CONSTANTS

Handling exceptions: handler

- Save hardware context (SAVE_ALL)
- call exception service routine
- Restore hardware context (RESTORE_ALL)



YOU HAVE TO IMPLEMENT IT

- Remove (if any) the exception hardware error code from the kernel stack
- Return to user (iret)


Handling exceptions: service routines

- Not managed in delivery 1
- Service routines will just print a message and waits forever
 - If an exception happens, and the kernel doesn't fix it, the program cannot continue

Interrupts

-clock
-keyboard

Handling interrupts: overview



Handling interrupts: initialization

- Interrupts initialization
 - Init IDT entry
 - 32: clock,
 - 33: keyboard
 - You have to decide when (in the code sequence) the system is prepared to receive interrupts in a safe way.

Handling interrupts: handler

- Similar to exceptions but:
 - No hardware error parameter in the kernel stack
 - It is necessary to notify the controller with the end of the interrupt management
 - It means that a new interrupt can be managed
 - End Of Interrupt (EOI)


```
movb $0x20, %al
outb %al, $0x20
```

Keyboard interrupt service routine

- Keyboard interrupt service routine

- Write character at the down-right side of the screen
- Access to keyboard data register

- `int inb (int port)`



- Scan code must be translated to character using the `char_map` table (`interrupt.c`)

Clock interrupt service routine

- Clock Service routine

- Write "minutes: seconds" at the top-right side of the screen
- You will need to implement helping functions
 - `Itoa`
 - `Printk_xy`
 - See the documentation

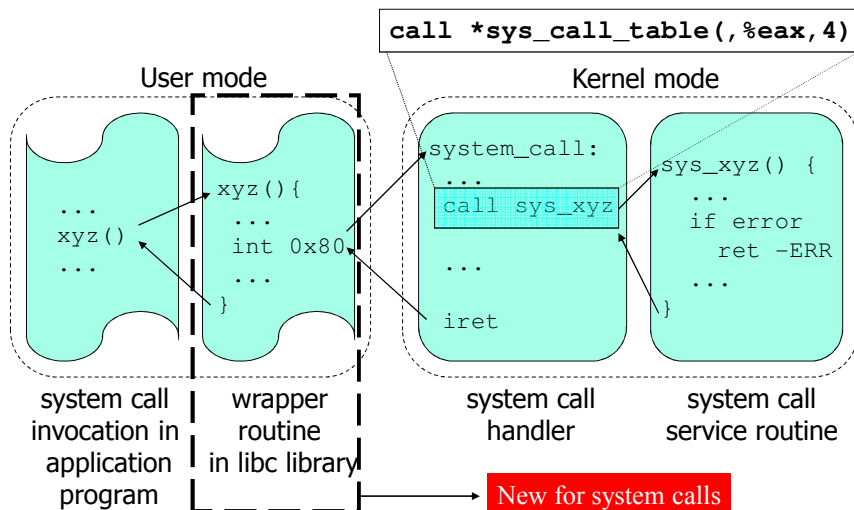
System calls

-write

Handling system calls

- System calls are generated by users
 - Intel: int assembly instruction (int idt_entry)
 - We must provide the users with an easy and portable way to invoke them
 - New layer: wrappers (libc.c)
 - User code is not reliable
 - **OS must validate all the data provided by the users**
 - Parameters
 - System calls identifiers
 - You have to decide which type of validation and where to include it (read the documentation)

Handling system calls: overview



Handling system calls: steps

- Initialization
- Wrapper routine
- System call handler
 - 1 for all the system calls
 - Redirects using a system call table
- System call service routine
 - 1 per system call

Handling system calls: initialization

- Init IDT
 - IDT entry: 128 (0x80)
 - void setTrapHandler (int n, void (*f)(), int DPL)
 - Similar to interrupts and exceptions

Handling system calls: wrapper (I)

- Written in assembly language and called by user C code
 - Observe the C compiler conventions
 - Which registers must be saved
 - How are parameters stacked
 - How results are returned
- Invoke the system call handler
 - Pass parameters
 - Identify the system call
 - Generate the trap
- Return the result to the user code

Handling system calls: wrapper (II)

- Pass parameters: Stack is not shared
 - Copy parameters from user stack to the registers
 - (first parameter) ebx, ecx, edx, esi, edi
 - Parameters in C are stacked from right to left
- Identify the system call
 - Copy system call number to eax
- Generate Trap: int \$0x80
- Return to user code
 - If error: returns -1 and updates the libc errno variable

Handling system calls: handler

- Save hardware context
 - Registers with system call parameters are located in the top of the kernel stack
- Execute system call service routine
 - Specified by eax (error checking)
 - Using system_call_table
 - call *sys_call_table(,%eax,4)
- Update kernel context with the system call result
- Restore context
- Return to user

PARAMETERS(which,where)?

Handling system calls: service routines

- System call service routine
 - Check parameters
 - Access the process address space (if needed)
 - `copy_from_user`
 - `copy_to_user`
 - Specific system call code
 - Can include the invocation to a device dependent routine
 - DELIVERY 2
 - `Sys_write_console`

PARAMETERS (which,where)?

TO DO

- Handling exceptions
- Handling interruptions: clock and keyboard
- Handling system calls: write
- Some additional functions
 - `Printk_xy`
 - `Itoa`
 - See documentation!

Deliverable P1.2

Process management

- Kernel Data structures
- Memory organization
- Process Scheduling



Summary

- Concepts
- Kernel data structures: process management
- Init process
- Memory organization
- System calls
- Scheduling
 - Context switch
 - Policies
- Process monitoring
- Process synchronization




Concepts

- Process

- “instance of a program in execution”
 - Understanding the Linux kernel. Pag. 72
- Resources are allocated to processes
- The process has a life cycle
- Processes are grouped in lists
 - Running, Waiting for I/O, Etc
- Hierarchical relationship
 - 1 father → N children

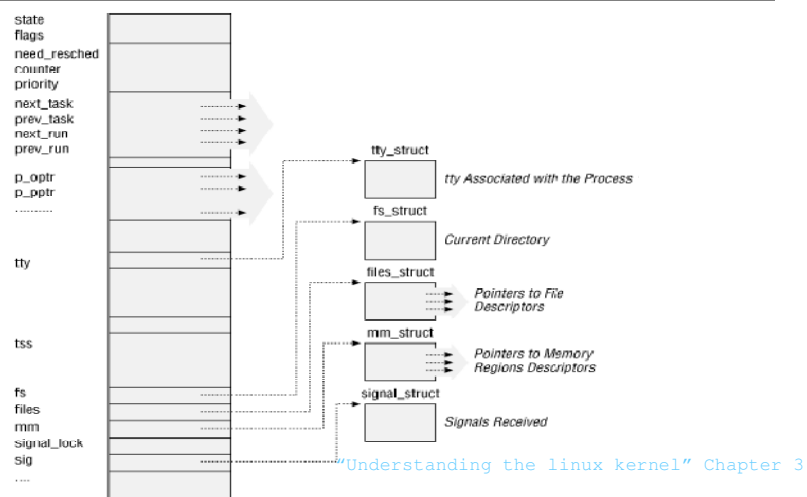
Concepts: Process life cycle



Concepts: process execution context

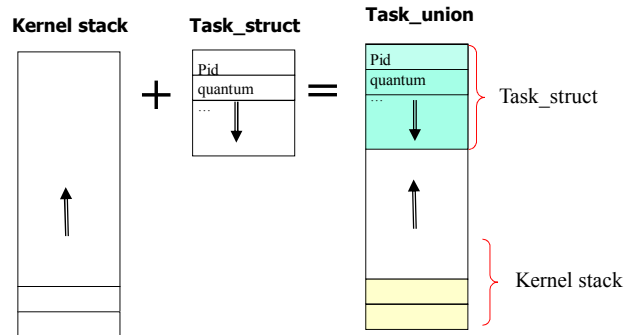
- The process execution context is:
 - The process address space : code,data,stack
 - The hardware context (registers value)
 - TSS (Task Segment Selector)
 - Kernel stack
- The hardware context is shared (1 set of registers). We must SAVE and RESTORE them
 - When changing the execution mode user \leftrightarrow kernel
 - When performing a context switch procA \leftrightarrow procB

Kernel data structures: task_struct



Kernel data structures: task union


```
union task_union {
    struct task_struct task;
    unsigned long stack[1024];};
```




Kernel data structures: Task's array

- Array 'task' contains all tasks.
- To detect stack overflows each task is protected by a special page.

```
struct protected_task_struct {
    unsigned long task_protect [KERNEL_STACK_SIZE];
    union task_union t;
};
struct protected_task_struct task [NR_TASKS];
```



Kernel data structures: lists



Kernel data structures

- How can we access the pointer to the running **task_struct** when entering the kernel?


The kernel stack and the **task_struct** **share** the same memory page (4k)



We know the kernel stack pointer of the running process (**esp**)



Based on the **esp** (applying a mask) we can obtain the address of the **task_struct** of the running process (**current**)



Kernel data structures

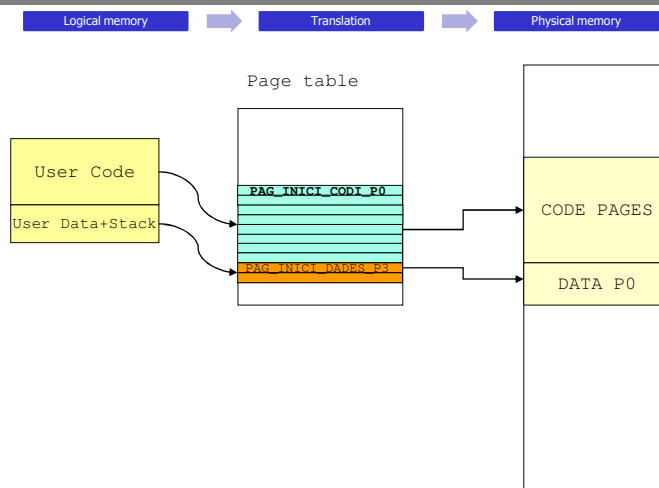
- Main work to do:
 - Complete type definitions
 - task_struct
 - Implement Kernel data needed
 - Task_struct array
 - Runqueue list

Init process

- Previously implemented
 - System/user Address space (see mm.c)
 - User/system stack allocated
 - TSS initialization (See setTSS function)
- You must:
 - Initialize task_struct for init process
 - PID=0
 - User stack
 - system stack
 - Insert it in runqueue list

Memory organization

From logical to Physical addresses



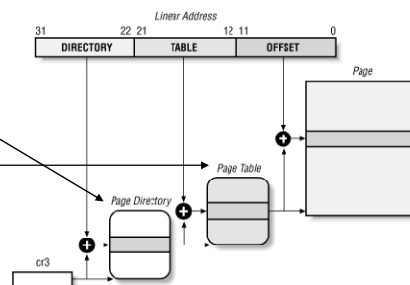
Segmentation+Pagination



- Segmentation
 - Initialized in Zeos and fixed → no modifications are required (see documentation P1)
- Paging Unit

dir_pagines in Zeos (mm.c)

taula_pagusr in Zeos (mm.c)

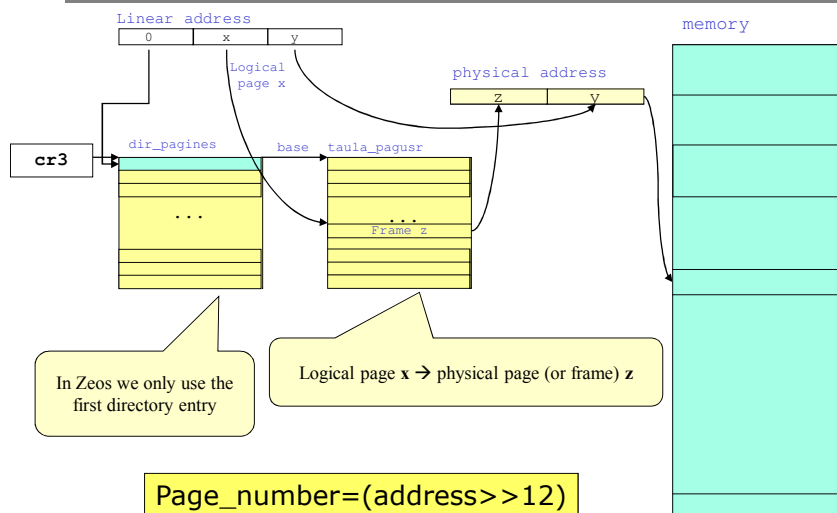


FIB

77



Organization in Zeos



FIB

78



Data to manage memory(mm.c)

```

/* SEGMENTATION */
/* Memory segments description table */
Descriptor *gdt = (Descriptor *) GDT_START;
/* Register pointing to the memory segments table */
Register gdtR;

/* PAGING */
/* Variables containing the page directory and the page table */

page_table_entry dir_pages[TOTAL_PAGES]
__attribute__((__section__(".data.task")));


page_table_entry pagusr_table[TOTAL_PAGES]
__attribute__((__section__(".data.task")));

/* TSS */
TSS tss;


```

Zeos Logical address space: consecutive

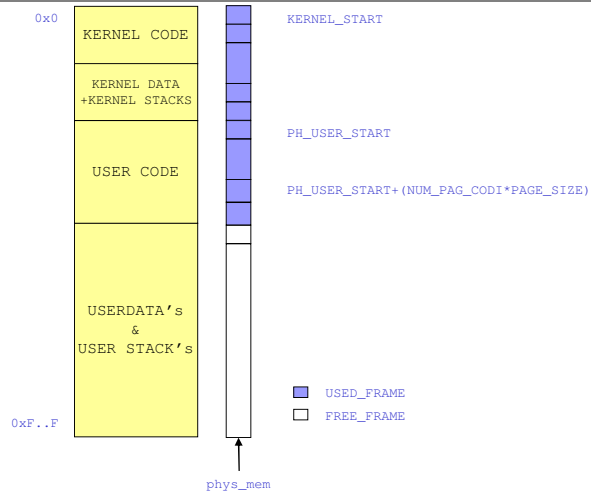
- User mode



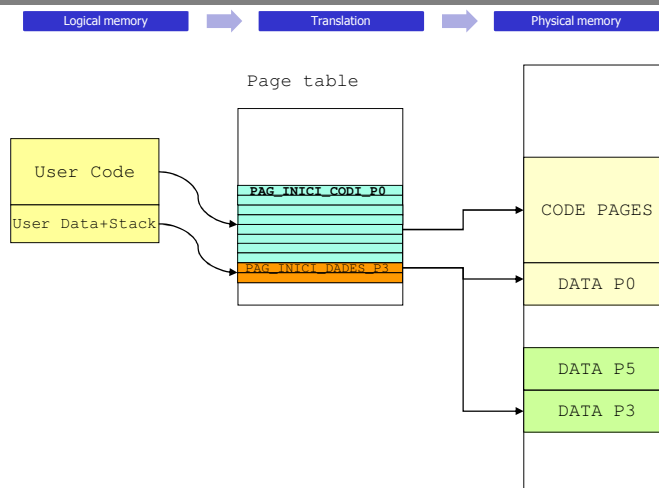
- Kernel mode



ZeOS Physical Memory Layout



Changing address space: P0 → P3



System calls

- getpid → returns process ID (PID)
- fork → creates a child process
- nice → modifies the process quantum
- exit → finalizes a process
- get_stats → returns monitoring info.

Fork: main steps

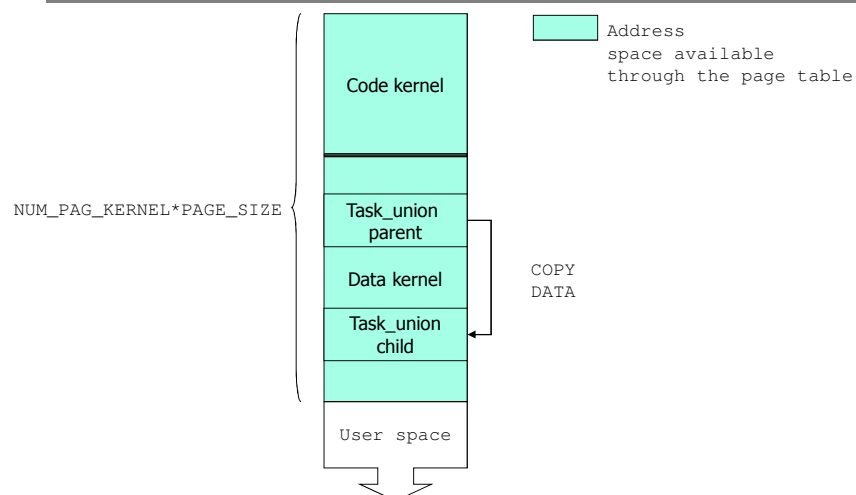
- Get free entry in task table
- Inherit system data
- Inherit user data
- Assign new PID (unique!)
- Update task_struct data
 - Includes I/O data
- Insert process in runqueue list
- Return pid of child process



Fork: Inherit system data

- 2-minutes to think about:
 1. Do we have access to the data?
 2. Are the data shared by both processes or not?
- System data: We have access to all the kernel data → no actions are required
 - Code → shared → no actions
 - Data → shared → no actions
 - Task_union (stack+task_struct) → **private** → we must copy from parent to child process


System data



Fork: Inherit user data

- The user address space is private
- If we have to copy some parts we have to modify the address space of the father
- **Code:** shared
 - no actions are required
- **Data+stack: private**
 - Allocate free frames for child data+stack
 - Grant access from parent to the child frames
 - **modify the memory page table**
 - Copy data+stack from parent to child
 - Deny access from parent to the child frames

User address space: Ex. P0 creates P5



Exit

- Update task_struct data
- Free allocated structures
 - Deallocate data+stack frames
 - Sempahores
 - ...
- Delete process for the runqueue list
- Free entry of tasks table
- *Schedule*

Page fault management

- An example of exceptions management
 - Kill the process that causes the exception
 - Similar to the exit system call
 - Show a message in the screen
 - Continue with the execution of the rest of the processes

Process Scheduling

- The policy:
 - Evaluates if a change must be performed
 - And selects the *next* process
- The kernel performs a *context switch*
- Goals:
 - Fast response time
 - Good throughput
 - Prevent process starvation
 - Priority management
 - etc.

Process Scheduling: Round Robin


- Each process receives a (per process) unit of execution: One quantum= N tics
- When the quantum expires or process blocks
 - New process is selected for execution
 - First process of runqueue is always selected
 - Before selection, current process must be dequeued and enqueued
- When a process is selected to run, a whole quantum is assigned to it
- You must:
 - Define & implement required scheduler functions
 - Define & implement required round robin functions

Context switch

```
Scheduling()
{
    Task next;
    UpdateSchedulingData();
    if (MustChangeProcess()){
        Next=SelectNextProcess();
        Context_switch(next);
    }
}
Context_switch(P)
{
    SAVE(CURRENT)
    RESTORE(P)
}
```

Context switch

- The kernel suspends the execution of the running process and resume the execution of some “previously” suspended process
 - SAVE the execution context of the running process and RESTORE the execution context of the suspended one
 - Address space
 - TSS
 - Kernel stack
 - Hardware context



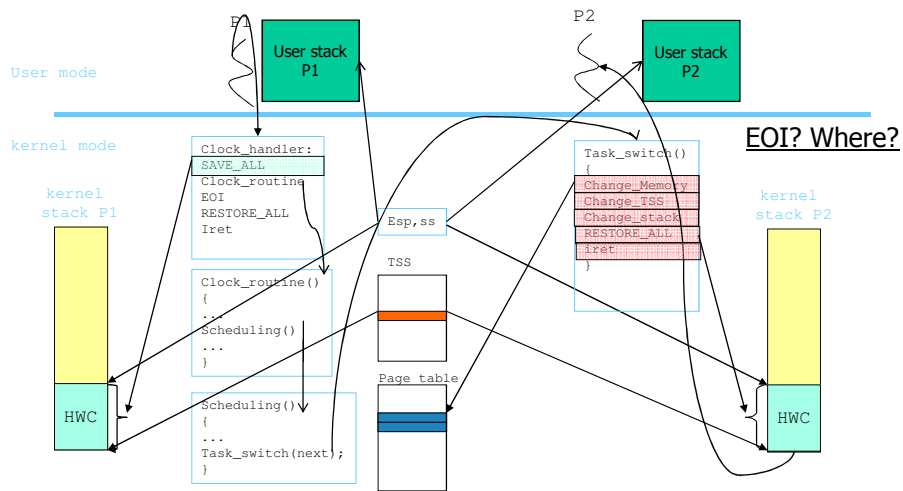
Context switch:SAVE

- Where/which is the context of current process saved?
 - Address space
 - Content: private, no need to save it
 - However, the info must be stored in task_struct
 - TSS (kernel stack address)
 - replaced each time, no need to save it
 - Address known, no need to save it
 - Kernel stack
 - Content:private, no need to save it
 - Address known, no need to save it
 - Hardware context
 - **It is saved in the kernel stack when entering the kernel (SAVE_ALL macro)**
 - **Where? Fixed position in kernel stack**

Context switch:RESTORE

- Resuming the context of *next* process
 - Address space
 - User (data+stack) pages
 - System has to move to the required page table entries
 - TSS
 - It must point to the kernel stack of *next* process (esp0)
 - Hardware context
 - esp registers → must point to the saved context of *next* process
 - The rest of registers are restored from the *next* kernel stack

Process context switch:example



Process monitoring

- **get_stats** → returns per process monitoring info.

```

struct stats {
int total_tics; /* Total tics executed by the process */
int total_trans; /* Total transitions ready → run */
int remaining_tics; /* Remaining tics to end the quantum */
};
  
```

```

int get_stats(int pid, struct stat *st)
  
```

Semaphores

Process Synchronization

- Semaphores
 - Creation: **sem_init**
 - Synchronization: **sem_wait** and **sem_signal**
 - Destruction: **sem_destroy**
 - 25-30 semaphores
- The use of semaphores can block and unblock processes (**NEW!**)

Creation

- `int sem_init(int n_sem, unsigned int value)`
 - `n_sem`: semaphore identifier
 - `value`: initial value for semaphore's counter
 - Returns -1 if error or 0 if ok
 - System call Id is 21
 - Initializes for semaphore '`n_sem`':
 - counter to '`value`'
 - blocked processes queue
 - Process that initializes the semaphore becomes the owner

`sem_wait()`

- `int sem_wait(int n_sem)`
 - `n_sem`: semaphore id
 - Returns: -1 if error or 0 if ok
 - System call Id is 22
 - If counter of semaphore '`n_sem`' $\leq 0 \rightarrow$
 - Block current process
 - Else
 - Decrement counter

sem_signal()

- **int sem_signal(int n_sem)**
 - n_sem: semaphore id
 - Returns: -1 if error or 0 if ok
 - System call Id is 23
 - If no blocked process at semaphore 'n_sem' →
 - Increment counter
 - Else
 - Unblock First blocked process

Destruction

- **int sem_destroy(int n_sem)**
 - n_sem: semaphore id
 - Returns: -1 if error or 0 if ok
 - Unblocks blocked processors (sem_wait will return -1) and releases the semaphore
 - System call Id is 24
 - Errors:
 - semaphore not initialized
 - Calling process not the owner

Deliverable P1.3

Input/Output management

- Logical devices
- Virtual devices
- zeosFAT File system

FIB

Delivery 1.3 overview

- Processes needs to have access to/from devices
- OS features
 - Hides physical characteristics
 - Provides new “logical” devices
 - Manages devices efficiently, securely
 - Offers virtual devices to processes

FIB

106



Delivery 1.3 overview

- Some information must be persistent
 - Stored in disk
- File systems organize and manage data stored in disks
 - zeosFAT: File system based on FAT
- You have to:
 - Design new data structures
 - Logical/virtual devices
 - Implement system calls for input/output
 - Design/Implement zeosFAT

Devices

- Physical: Hardware devices
- Logical
 - known by the O.S
 - OS internal data representation + Device dependent functions
 - Abstract zero, one, or multiple physical devices
- Virtual (file descriptors)
 - Known by processes
 - Managed through system calls

Physical devices

- Keyboard
 - Read only
- Display
 - Write only

Logical devices

- Represents zero, one or multiple devices
 - Zero: ex. Pipes
 - One: keyboard
 - Multiple: console (keyboard+display)
- The OS defines a set of functions to make uniform the device access
 - New data type: file_operations
 - Each device can have its own data
- Files will represent logical devices

Logical devices: file_operations

- Define the common set of functions to access to devices
 - Open
 - Read
 - Write
 - Etc
- Same API to all the devices
 - The set of functions and parameters of each one are a superset

Logical devices: files

- Static characteristics
 - Name (to simplify, no multiples names for one logic device is allowed)
 - Access mode allowed: read, write, read&write
 - File_operations
- Dynamic characteristics
 - Depends on the specific accesses in course
 - Current offset
 - Open mode
 - ...

Logical devices: files

- Define the required OS structures to guarantee (see system calls semantics):
 - Two processes accessing concurrently the same file
 - Concurrent but independent
 - Two processes sharing the access to the same file
 - Sharing dynamic data
 - One process accessing concurrently the same file multiple times
 - One process sharing the access to the same file multiple times

Logical devices: zeos name space

- Zeos specific: not persistent, only in memory
- Single level of directory
 - Not sub-directories are allowed
- Design the directory data type and the required functionality to initialize it at the start of zeos.

Virtual devices

- Each time a process ask for access of a file (open's it), the OS returns a new per-process handler (file descriptor)
- The handler binds process, the logical device, and the dynamic characteristics associated with the open
- Default fd
 - stdin : fd=0
 - stdout: fd=1
 - stderr: fd=2

Virtual devices

- Each process has its own table of file descriptors
- When a process calls forks, the child process heritages the file descriptor table (a copy)
 - fd's of child process will SHARE dynamic information with the parent process

New System calls

- Open
- Read
- Dup
- Close

Syscall open

- `int open(char *path, int mode)`
 - path: name of the file to be opened.
 - mode can be:
 - `O_RDONLY` (0x1)
 - `O_WRONLY` (0x2)
 - `O_RDWR` (0x3)
 - Returns: -1 if error or a new fd if ok
 - System call Id is 5
 - Creates an open file descriptor that refers to the file and establishes a connection between it and the returned fd.

Syscall read

- `int read(int fd, char *buff, int size)`
 - fd: file descriptor to read from
 - buff: pointer where data will be copied to
 - size: size (in bytes) to read
 - Return: -1 if error or number of read bytes if OK
 - System call Id is 3

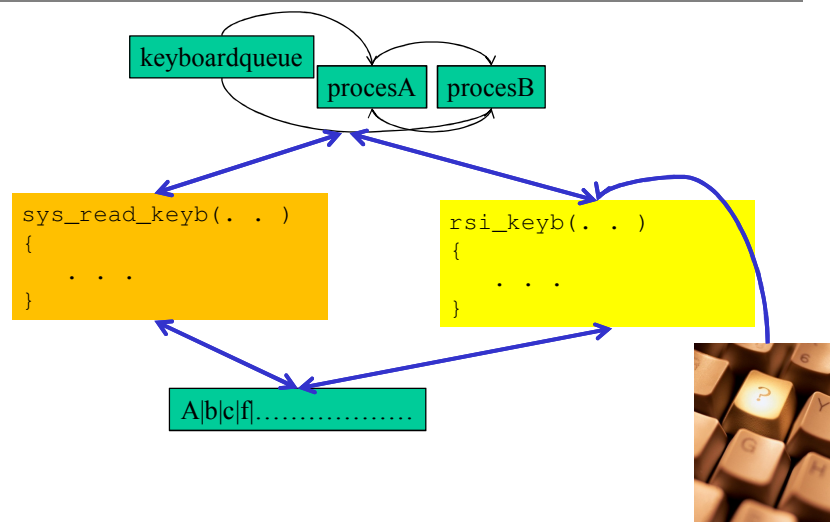
Note that:

- Buff must be in the user address space
- Read's are served in arrival order
- Calls the device dependent function

Syscall: dup/close

- `int dup (int fd)`
 - fd: file descriptor to duplicate
 - Return: -1 if error or the first fd available if OK
 - System call Id is 41
 - Duplicate an open file descriptor. The new fd SHARES the dynamic information
- `int close (int fd)`
 - fd: file descriptor to close
 - Return: -1 if error or 0 if OK
 - System call Id is 6
 - Unbinds file descriptor and its logical device

Reading from the keyboard



Sys_read_keyboard

- Shared buffer between this function and the keyboard rsi
 - Circular buffer (new type)
- If there are pending requests
 - BLOCK(keyboardqueue) the calling process
- Otherwise
 - If there are enough bytes to read → read and return
 - Otherwise → BLOCK(keyboardqueue) the process
- Before blocking, we need to save the "request" data (which info?)

Rsi Keyboard

- Stores new characters in the shared buffer
- If there are blocked processes in the keyboardqueue
 - If there are enough bytes to finish the request OR the buffer is full
 - Copy data → Be careful, process A, executing the rsi, is copying data to user space of process B

ZeosFAT

File System

- Defines how files are stored in disk and how the disk space is managed
 - Allocates blocks
 - Organize blocks
 - Etc
- Per file system you have to decide
 - How to allocate new blocks
 - How to manage free blocks

zeosFAT

- In memory file system, we are not going to access to disk (not persistent)
- We will substitute the disk by a vectors of data blocks of 256 bytes
 - New data type
 - Design it and think about functionality

How to allocate disk space?

- The unit of allocation will be blocks
- Blocks in a file will be linked
 - Each block will point to the next of the same file
- Pointers to next block will be separated from data
 - New data type: File Allocation Table (FAT)
- We need to know the first block of data per file

How to manage free blocks?

- Using the same FAT
 - Free blocks are linked
 - We need a pointer to the first free block

Directory

- Single level of directory

Existing System calls

- open
 - Add new flags to create a file:
 - O_CREAT (0x4)
 - O_EXCL (0x8)
- write/read/dup/close
 - Should remain without modifications

System call to add

- `int unlink(const char *path)`
 - Path: name of the file to be removed
 - Return: -1 if error (not abled to remove) or 0 if OK
 - System call Id is 10
 - Remove a file. A file only can be removed if no processes are using it: you should maintain the number of active references

zeosFAT

- Design the new data types